

3_Modeling_V2

May 3, 2021

```
[84]: import modin.pandas as pd
import numpy as np

import pickle

from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, plot_confusion_matrix

import keras
from keras.models import Sequential
from keras.layers import Dense, Dropout
import keras.backend as K
from keras.utils import plot_model

import matplotlib.pyplot as plt
import seaborn as sns

import shap
```

0.1 Loading training data

```
[2]: training_dat = pd.read_pickle('./PROCESSED/training_dat2.pkl')
training_dat.columns
```

UserWarning: Ray execution environment not yet initialized. Initializing..
To remove this warning, run the following python code before doing dataframe operations:

```
import ray
ray.init()
```

UserWarning: `read_pickle` defaulting to pandas implementation.
To request implementation, send an email to feature_requests@modin.org.

```
[2]: Index(['Loan Sequence Number', 'Loan Delinquency Within Year', 'Credit Score',
'MortgageInsuranceFlag', 'Units_1', 'Units_2', 'Units_3', 'Units_4',
'OccupancyStatus_I', 'OccupancyStatus_P', 'OccupancyStatus_S',
```

```
'Original Combined Loan-to-Value (CLTV)',
'Original Debt-to-Income (DTI) Ratio', 'Original UPB',
'Original Loan-to-Value (LTV)', 'Original Interest Rate', 'Channel_B',
'Channel_C', 'Channel_R', 'PropertyType_CO', 'PropertyType_CP',
'PropertyType_MH', 'PropertyType_PU', 'PropertyType_SF',
'LoanPurpose_C', 'LoanPurpose_N', 'LoanPurpose_P', 'LoanTerm_360',
'LoanTerm_180', 'LoanTerm_240', 'LoanTerm_120', 'LoanTerm_300',
'Original Loan Term', 'OneBorrower', 'AffordableProgramFlag'],
dtype='object')
```

```
[3]: len(training_dat)
```

```
[3]: 396267
```

```
[4]: training_dat['Loan Delinquency Within Year'].value_counts()
```

UserWarning: value_counts defaulting to pandas implementation.

```
[4]: 0.0    387832
     1.0    8435
     Name: Loan Delinquency Within Year, dtype: int64
```

0.2 Selecting features to feed model

Currently not excluding any columns based on the metrics. I have manually excluded a couple.

```
[6]: metrics = training_dat.describe().transpose()
     metrics.reset_index(inplace=True)
     metrics.rename(columns={'index': 'column'}, inplace=True)
     metrics.head()
```

```
[6]:
```

	column	count	mean	std	min	\
0	Loan Delinquency Within Year	396267.0	0.021286	0.144337	0.000000	
1	Credit Score	396267.0	0.000578	0.999393	-5.199338	
2	MortgageInsuranceFlag	396267.0	0.218746	0.413397	0.000000	
3	Units_1	396267.0	0.980379	0.138693	0.000000	
4	Units_2	396267.0	0.013582	0.115747	0.000000	

	25%	50%	75%	max
0	0.000000	0.000000	0.000000	1.000000
1	-0.672129	-0.003764	0.670557	5.199338
2	0.000000	0.000000	0.000000	1.000000
3	1.000000	1.000000	1.000000	1.000000
4	0.000000	0.000000	0.000000	1.000000

```
[7]: tf1 = metrics['mean'] < 1.1
     tf2 = metrics['mean'] > -.1
```

```
tf = tf1 & tf2
metrics[tf]
```

```
[7]:
```

	column	count	mean	std \
0	Loan Delinquency Within Year	396267.0	0.021286	0.144337
1	Credit Score	396267.0	0.000578	0.999393
2	MortgageInsuranceFlag	396267.0	0.218746	0.413397
3	Units_1	396267.0	0.980379	0.138693
4	Units_2	396267.0	0.013582	0.115747
5	Units_3	396267.0	0.003147	0.056009
6	Units_4	396267.0	0.002892	0.053699
7	OccupancyStatus_I	396267.0	0.072335	0.259042
8	OccupancyStatus_P	396267.0	0.886062	0.317737
9	OccupancyStatus_S	396267.0	0.041603	0.199681
10	Original Combined Loan-to-Value (CLTV)	396267.0	-0.007684	0.972985
11	Original Debt-to-Income (DTI) Ratio	396267.0	-0.004331	0.997126
12	Original UPB	396267.0	0.000987	0.997631
13	Original Loan-to-Value (LTV)	396267.0	-0.013095	0.972578
14	Original Interest Rate	396267.0	-0.003350	0.995451
15	Channel_B	396267.0	0.108051	0.310445
16	Channel_C	396267.0	0.316756	0.465212
17	Channel_R	396267.0	0.575193	0.494314
18	PropertyType_CO	396267.0	0.068686	0.252920
19	PropertyType_CP	396267.0	0.001340	0.036582
20	PropertyType_MH	396267.0	0.003455	0.058676
21	PropertyType_PU	396267.0	0.253067	0.434769
22	PropertyType_SF	396267.0	0.673452	0.468951
23	LoanPurpose_C	396267.0	0.240255	0.427239
24	LoanPurpose_N	396267.0	0.288422	0.453029
25	LoanPurpose_P	396267.0	0.471324	0.499178
26	LoanTerm_360	396267.0	0.737490	0.439999
27	LoanTerm_180	396267.0	0.187578	0.390375
28	LoanTerm_240	396267.0	0.047826	0.213399
29	LoanTerm_120	396267.0	0.015727	0.124417
30	LoanTerm_300	396267.0	0.006127	0.078036
31	Original Loan Term	396267.0	-0.064882	0.698888
32	OneBorrower	396267.0	0.461608	0.498525
33	AffordableProgramFlag	396267.0	0.008272	0.090575

	min	25%	50%	75%	max
0	0.000000	0.000000	0.000000	0.000000	1.000000
1	-5.199338	-0.672129	-0.003764	0.670557	5.199338
2	0.000000	0.000000	0.000000	0.000000	1.000000
3	0.000000	1.000000	1.000000	1.000000	1.000000
4	0.000000	0.000000	0.000000	0.000000	1.000000
5	0.000000	0.000000	0.000000	0.000000	1.000000
6	0.000000	0.000000	0.000000	0.000000	1.000000

7	0.000000	0.000000	0.000000	0.000000	1.000000
8	0.000000	1.000000	1.000000	1.000000	1.000000
9	0.000000	0.000000	0.000000	0.000000	1.000000
10	-5.199338	-0.683178	0.012546	0.376283	5.199338
11	-5.199338	-0.686350	-0.023839	0.616541	5.199338
12	-5.199338	-0.676854	0.010037	0.678433	5.199338
13	-5.199338	-0.686350	0.011291	0.421111	5.199338
14	-5.199338	-0.705530	-0.089192	0.716839	5.199338
15	0.000000	0.000000	0.000000	0.000000	1.000000
16	0.000000	0.000000	0.000000	1.000000	1.000000
17	0.000000	0.000000	1.000000	1.000000	1.000000
18	0.000000	0.000000	0.000000	0.000000	1.000000
19	0.000000	0.000000	0.000000	0.000000	1.000000
20	0.000000	0.000000	0.000000	0.000000	1.000000
21	0.000000	0.000000	0.000000	1.000000	1.000000
22	0.000000	0.000000	1.000000	1.000000	1.000000
23	0.000000	0.000000	0.000000	0.000000	1.000000
24	0.000000	0.000000	0.000000	1.000000	1.000000
25	0.000000	0.000000	0.000000	1.000000	1.000000
26	0.000000	0.000000	1.000000	1.000000	1.000000
27	0.000000	0.000000	0.000000	0.000000	1.000000
28	0.000000	0.000000	0.000000	0.000000	1.000000
29	0.000000	0.000000	0.000000	0.000000	1.000000
30	0.000000	0.000000	0.000000	0.000000	1.000000
31	-5.199338	-0.736442	0.334851	0.334851	5.199338
32	0.000000	0.000000	0.000000	1.000000	1.000000
33	0.000000	0.000000	0.000000	0.000000	1.000000

```
[8]: keep_cols = ['Loan Sequence Number']
exclude_cols = ['Original Combined Loan-to-Value',
                '(CLTV)', 'PropertyType_CP', 'Mortgage Insurance Percentage (MI %)', 'Original',
                'Loan Term']
good_cols = [col for col in list(metrics[tf]['column']) if col not in
              exclude_cols]
keep_cols.extend(good_cols)
keep_cols
```

```
[8]: ['Loan Sequence Number',
      'Loan Delinquency Within Year',
      'Credit Score',
      'MortgageInsuranceFlag',
      'Units_1',
      'Units_2',
      'Units_3',
      'Units_4',
      'OccupancyStatus_I',
      'OccupancyStatus_P',
```

```

'OccupancyStatus_S',
'Original Debt-to-Income (DTI) Ratio',
'Original UPB',
'Original Loan-to-Value (LTV)',
'Original Interest Rate',
'Channel_B',
'Channel_C',
'Channel_R',
'PropertyType_CO',
'PropertyType_MH',
'PropertyType_PU',
'PropertyType_SF',
'LoanPurpose_C',
'LoanPurpose_N',
'LoanPurpose_P',
'LoanTerm_360',
'LoanTerm_180',
'LoanTerm_240',
'LoanTerm_120',
'LoanTerm_300',
'OneBorrower',
'AffordableProgramFlag']

```

```
[9]: len(keep_cols)
```

```
[9]: 32
```

```
[10]: final_df = training_dat[keep_cols].copy()
final_df.head()
```

```
[10]:
```

	Loan Sequence Number	Loan Delinquency Within Year	Credit Score	\
0	F110Q1000008	0.0	1.315958	
1	F110Q1000064	0.0	2.225823	
2	F110Q1000072	0.0	0.451469	
3	F110Q1000080	0.0	-1.733071	
4	F110Q1000096	0.0	1.033647	

	MortgageInsuranceFlag	Units_1	Units_2	Units_3	Units_4	\
0	0	1	0	0	0	
1	0	1	0	0	0	
2	0	1	0	0	0	
3	0	1	0	0	0	
4	0	1	0	0	0	

	OccupancyStatus_I	OccupancyStatus_P	...	LoanPurpose_C	LoanPurpose_N	\
0	0	1	...	0	1	
1	0	1	...	1	0	

2	0	1	...	0	1
3	0	1	...	0	1
4	0	1	...	0	1

	LoanPurpose_P	LoanTerm_360	LoanTerm_180	LoanTerm_240	LoanTerm_120	\
0	0	1	0	0	0	
1	0	1	0	0	0	
2	0	1	0	0	0	
3	0	1	0	0	0	
4	0	1	0	0	0	

	LoanTerm_300	OneBorrower	AffordableProgramFlag
0	0	0	0
1	0	1	0
2	0	0	0
3	0	1	0
4	0	1	0

[5 rows x 32 columns]

```
[11]: len(final_df)
```

```
[11]: 396267
```

0.3 Splitting in to training and test datasets

```
[12]: idx_col = 'Loan Sequence Number'

test_size = round(len(final_df[idx_col]) * .2)
test_loans = final_df[idx_col].sample(test_size)
train_loans = list(set(final_df[idx_col].unique()) - set(test_loans))

train_df = pd.DataFrame({idx_col:train_loans})
train_df = train_df.merge(final_df, on=idx_col, how='left')

test_df = pd.DataFrame({idx_col:test_loans})
test_df = test_df.merge(final_df, on=idx_col, how='left')
```

UserWarning: Distributing <class 'dict'> object. This may take some time.

```
[13]: train_df.head()
```

```
[13]:   Loan Sequence Number  Loan Delinquency Within Year  Credit Score  \
0      F110Q1252148                0.0      -0.177827
1      F111Q4171377                0.0      -0.383024
2      F114Q3148996                0.0       0.634851
3      F118Q2045353                0.0      -0.415634
```

4	F114Q4174402		0.0	0.634851
---	--------------	--	-----	----------

	MortgageInsuranceFlag	Units_1	Units_2	Units_3	Units_4	\
0	0	1	0	0	0	
1	0	1	0	0	0	
2	0	1	0	0	0	
3	0	1	0	0	0	
4	1	1	0	0	0	

	OccupancyStatus_I	OccupancyStatus_P	...	LoanPurpose_C	LoanPurpose_N	\
0	0	1	...	1	0	
1	0	1	...	0	1	
2	0	1	...	1	0	
3	0	1	...	0	0	
4	0	1	...	0	0	

	LoanPurpose_P	LoanTerm_360	LoanTerm_180	LoanTerm_240	LoanTerm_120	\
0	0	0	1	0	0	
1	0	0	1	0	0	
2	0	0	1	0	0	
3	1	0	1	0	0	
4	1	1	0	0	0	

	LoanTerm_300	OneBorrower	AffordableProgramFlag
0	0	1	0
1	0	0	0
2	0	1	0
3	0	0	0
4	0	1	0

[5 rows x 32 columns]

```
[14]: test_df.head()
```

```
[14]:  Loan Sequence Number  Loan Delinquency Within Year  Credit Score  \
0      F117Q3099461      0.0      0.604448
1      F115Q1238542      0.0     -0.415634
2      F119Q2306763      0.0     -0.231641
3      F110Q2117701      0.0     -0.705530
4      F116Q1069514      0.0     -0.605955
```

	MortgageInsuranceFlag	Units_1	Units_2	Units_3	Units_4	\
0	0	1	0	0	0	
1	1	1	0	0	0	
2	0	1	0	0	0	
3	0	1	0	0	0	
4	0	1	0	0	0	

	OccupancyStatus_I	OccupancyStatus_P	...	LoanPurpose_C	LoanPurpose_N	\
0	0	1	...	1	0	
1	0	1	...	0	0	
2	0	1	...	1	0	
3	0	1	...	1	0	
4	0	1	...	0	0	

	LoanPurpose_P	LoanTerm_360	LoanTerm_180	LoanTerm_240	LoanTerm_120	\
0	0	1	0	0	0	
1	1	1	0	0	0	
2	0	1	0	0	0	
3	0	1	0	0	0	
4	1	1	0	0	0	

	LoanTerm_300	OneBorrower	AffordableProgramFlag
0	0	0	0
1	0	1	0
2	0	1	0
3	0	0	0
4	0	1	0

[5 rows x 32 columns]

```
[15]: # Build training and test datasets
X_train = train_df[train_df.columns[2:]].to_numpy()
y_train = train_df["Loan Delinquency Within Year"].to_numpy()

X_test = test_df[test_df.columns[2:]].to_numpy()
y_test = test_df["Loan Delinquency Within Year"].to_numpy()
```

```
[16]: X_train.shape
```

```
[16]: (317014, 30)
```

0.4 Building model

```
[17]: # Based on https://keras.io/examples/structured_data/imbalanced_classification/

pos_count = final_df['Loan Delinquency Within Year'].sum()

weight_for_0 = 1.0 / (len(final_df)-pos_count)
weight_for_1 = 1.0 / pos_count

model = Sequential()
model.add(Dense(30,activation='relu',input_shape=(X_train.shape[-1],)))
model.add(Dense(30,activation='relu'))
```



```

model.add(Dropout(.3))
model.add(Dense(15,activation='relu'))
model.add(Dense(1,activation='sigmoid'))

metrics = [
    keras.metrics.FalseNegatives(name="fn"),
    keras.metrics.FalsePositives(name="fp"),
    keras.metrics.TrueNegatives(name="tn"),
    keras.metrics.TruePositives(name="tp"),
    keras.metrics.BinaryAccuracy(name="binary_accuracy"),
    keras.metrics.Precision(name="precision"),
    keras.metrics.Recall(name="recall")
]
model.compile(optimizer=keras.optimizers.Adam(1e-2),
    ↪loss="binary_crossentropy", metrics=metrics)

class_weight = {0: weight_for_0, 1: weight_for_1}

```

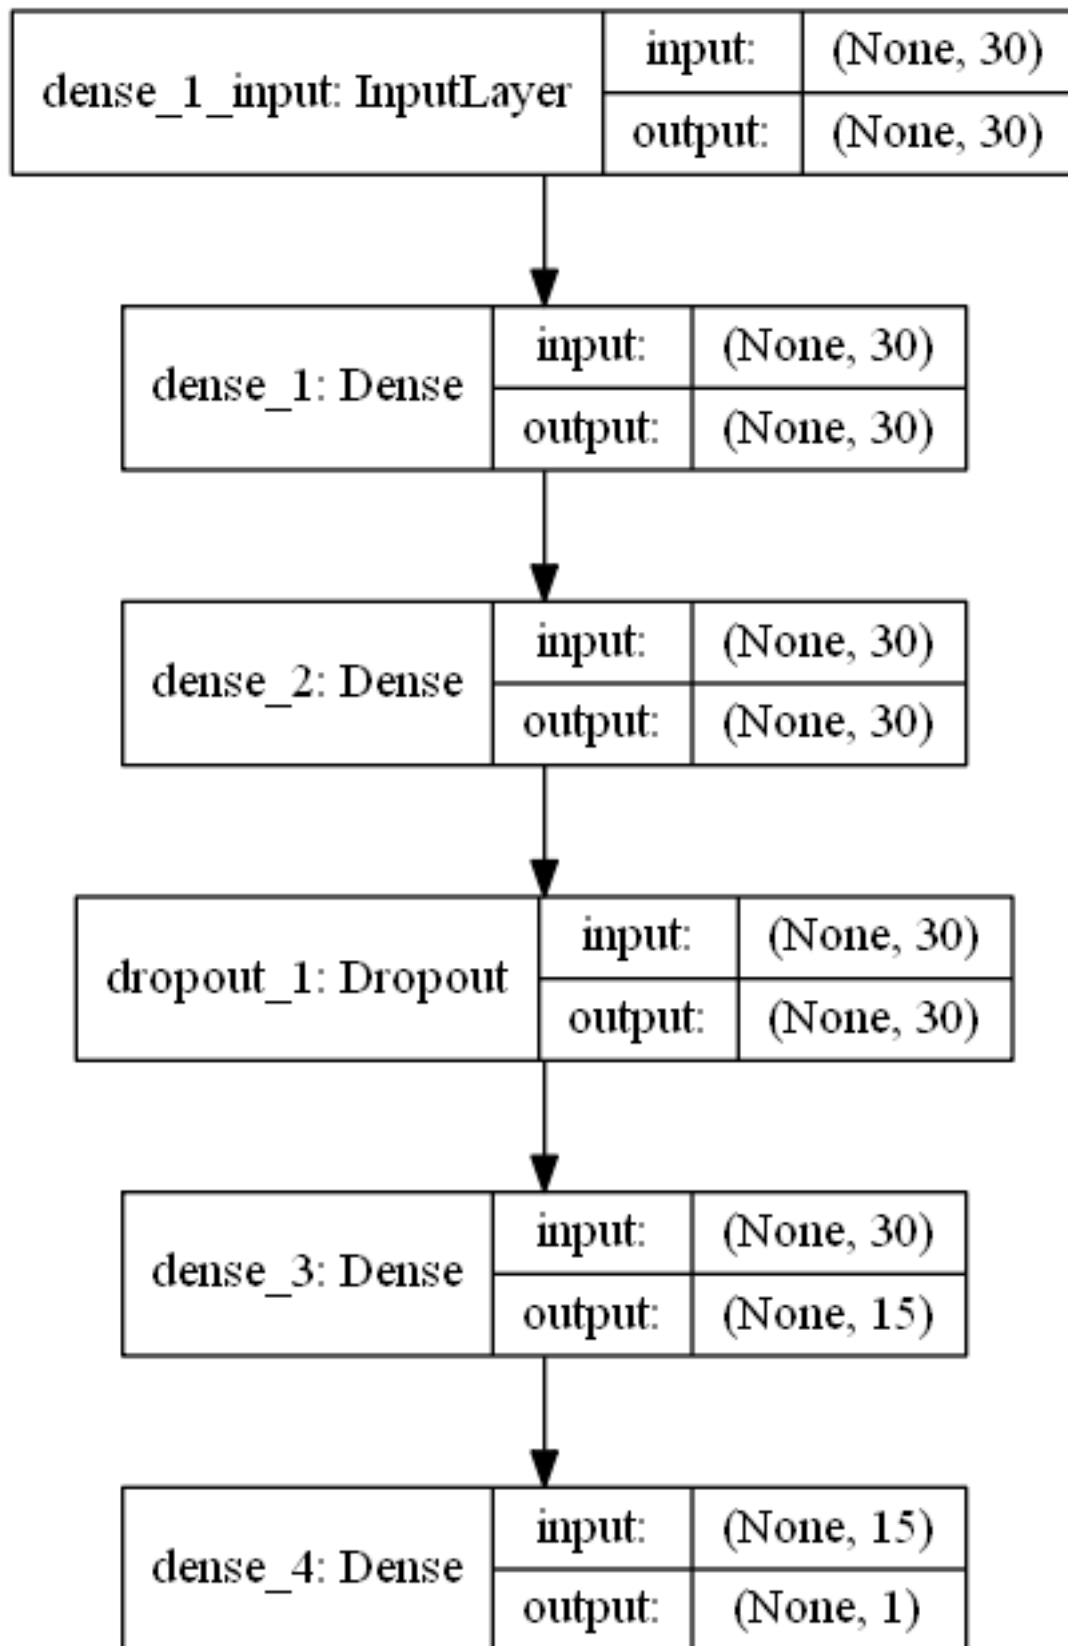
WARNING:tensorflow:From C:\tools\Anaconda3\lib\site-packages\keras\backend\tensorflow_backend.py:3172: add_dispatch_support.<locals>.wrapper (from tensorflow.python.ops.array_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

```
[18]: plot_model(model, show_shapes=True)
```

```
[18]:
```



0.5 Training model

```
[19]: history = model.fit(X_train,y_train,class_weight=class_weight,batch_size=2048,
                           epochs=20, validation_split=.3, verbose=1)
```

```
WARNING:tensorflow:From C:\tools\Anaconda3\lib\site-
packages\keras\backend\tensorflow_backend.py:422: The name tf.global_variables
is deprecated. Please use tf.compat.v1.global_variables instead.
```

Train on 221909 samples, validate on 95105 samples

Epoch 1/20

```
221909/221909 [=====] - 1s 5us/step - loss: 3.5002e-06
- fn: 1266.0000 - fp: 142513.0000 - tn: 74591.0000 - tp: 3539.0000 -
binary_accuracy: 0.3521 - precision: 0.0242 - recall: 0.7365 - val_loss:
3.2969e-06 - val_fn: 771.0000 - val_fp: 34768.0000 - val_tn: 58344.0000 -
val_tp: 1222.0000 - val_binary_accuracy: 0.6263 - val_precision: 0.0340 -
val_recall: 0.6131
```

Epoch 2/20

```
221909/221909 [=====] - 1s 3us/step - loss: 3.2980e-06
- fn: 1926.0000 - fp: 72892.0000 - tn: 144212.0000 - tp: 2879.0000 -
binary_accuracy: 0.6628 - precision: 0.0380 - recall: 0.5992 - val_loss:
3.1935e-06 - val_fn: 869.0000 - val_fp: 26800.0000 - val_tn: 66312.0000 -
val_tp: 1124.0000 - val_binary_accuracy: 0.7091 - val_precision: 0.0403 -
val_recall: 0.5640
```

Epoch 3/20

```
221909/221909 [=====] - 1s 3us/step - loss: 3.2257e-06
- fn: 1896.0000 - fp: 68305.0000 - tn: 148799.0000 - tp: 2909.0000 -
binary_accuracy: 0.6836 - precision: 0.0408 - recall: 0.6054 - val_loss:
3.1776e-06 - val_fn: 901.0000 - val_fp: 25340.0000 - val_tn: 67772.0000 -
val_tp: 1092.0000 - val_binary_accuracy: 0.7241 - val_precision: 0.0413 -
val_recall: 0.5479
```

Epoch 4/20

```
221909/221909 [=====] - 1s 3us/step - loss: 3.2080e-06
- fn: 1800.0000 - fp: 69277.0000 - tn: 147827.0000 - tp: 3005.0000 -
binary_accuracy: 0.6797 - precision: 0.0416 - recall: 0.6254 - val_loss:
3.1605e-06 - val_fn: 827.0000 - val_fp: 28218.0000 - val_tn: 64894.0000 -
val_tp: 1166.0000 - val_binary_accuracy: 0.6946 - val_precision: 0.0397 -
val_recall: 0.5850
```

Epoch 5/20

```
221909/221909 [=====] - 1s 4us/step - loss: 3.1852e-06
- fn: 1786.0000 - fp: 69270.0000 - tn: 147834.0000 - tp: 3019.0000 -
binary_accuracy: 0.6798 - precision: 0.0418 - recall: 0.6283 - val_loss:
3.1568e-06 - val_fn: 673.0000 - val_fp: 34625.0000 - val_tn: 58487.0000 -
val_tp: 1320.0000 - val_binary_accuracy: 0.6289 - val_precision: 0.0367 -
val_recall: 0.6623
```

Epoch 6/20

221909/221909 [=====] - 1s 4us/step - loss: 3.1806e-06
- fn: 1666.0000 - fp: 73660.0000 - tn: 143444.0000 - tp: 3139.0000 -
binary_accuracy: 0.6606 - precision: 0.0409 - recall: 0.6533 - val_loss:
3.1467e-06 - val_fn: 729.0000 - val_fp: 31564.0000 - val_tn: 61548.0000 -
val_tp: 1264.0000 - val_binary_accuracy: 0.6604 - val_precision: 0.0385 -
val_recall: 0.6342

Epoch 7/20

221909/221909 [=====] - 1s 3us/step - loss: 3.1780e-06
- fn: 1698.0000 - fp: 72418.0000 - tn: 144686.0000 - tp: 3107.0000 -
binary_accuracy: 0.6660 - precision: 0.0411 - recall: 0.6466 - val_loss:
3.1462e-06 - val_fn: 748.0000 - val_fp: 30406.0000 - val_tn: 62706.0000 -
val_tp: 1245.0000 - val_binary_accuracy: 0.6724 - val_precision: 0.0393 -
val_recall: 0.6247

Epoch 8/20

221909/221909 [=====] - 1s 3us/step - loss: 3.1696e-06
- fn: 1680.0000 - fp: 72347.0000 - tn: 144757.0000 - tp: 3125.0000 -
binary_accuracy: 0.6664 - precision: 0.0414 - recall: 0.6504 - val_loss:
3.1552e-06 - val_fn: 894.0000 - val_fp: 24494.0000 - val_tn: 68618.0000 -
val_tp: 1099.0000 - val_binary_accuracy: 0.7331 - val_precision: 0.0429 -
val_recall: 0.5514

Epoch 9/20

221909/221909 [=====] - 1s 4us/step - loss: 3.1630e-06
- fn: 1756.0000 - fp: 69704.0000 - tn: 147400.0000 - tp: 3049.0000 -
binary_accuracy: 0.6780 - precision: 0.0419 - recall: 0.6345 - val_loss:
3.1487e-06 - val_fn: 855.0000 - val_fp: 26170.0000 - val_tn: 66942.0000 -
val_tp: 1138.0000 - val_binary_accuracy: 0.7158 - val_precision: 0.0417 -
val_recall: 0.5710

Epoch 10/20

221909/221909 [=====] - 1s 4us/step - loss: 3.1581e-06
- fn: 1685.0000 - fp: 72752.0000 - tn: 144352.0000 - tp: 3120.0000 -
binary_accuracy: 0.6646 - precision: 0.0411 - recall: 0.6493 - val_loss:
3.1425e-06 - val_fn: 807.0000 - val_fp: 27579.0000 - val_tn: 65533.0000 -
val_tp: 1186.0000 - val_binary_accuracy: 0.7015 - val_precision: 0.0412 -
val_recall: 0.5951

Epoch 11/20

221909/221909 [=====] - 1s 4us/step - loss: 3.1548e-06
- fn: 1732.0000 - fp: 71073.0000 - tn: 146031.0000 - tp: 3073.0000 -
binary_accuracy: 0.6719 - precision: 0.0414 - recall: 0.6395 - val_loss:
3.1622e-06 - val_fn: 896.0000 - val_fp: 24234.0000 - val_tn: 68878.0000 -
val_tp: 1097.0000 - val_binary_accuracy: 0.7358 - val_precision: 0.0433 -
val_recall: 0.5504

Epoch 12/20

221909/221909 [=====] - 1s 3us/step - loss: 3.1561e-06
- fn: 1624.0000 - fp: 74023.0000 - tn: 143081.0000 - tp: 3181.0000 -
binary_accuracy: 0.6591 - precision: 0.0412 - recall: 0.6620 - val_loss:
3.1366e-06 - val_fn: 744.0000 - val_fp: 30436.0000 - val_tn: 62676.0000 -
val_tp: 1249.0000 - val_binary_accuracy: 0.6722 - val_precision: 0.0394 -
val_recall: 0.6267

Epoch 13/20

221909/221909 [=====] - 1s 4us/step - loss: 3.1451e-06
- fn: 1632.0000 - fp: 73668.0000 - tn: 143436.0000 - tp: 3173.0000 -
binary_accuracy: 0.6607 - precision: 0.0413 - recall: 0.6604 - val_loss:
3.1423e-06 - val_fn: 766.0000 - val_fp: 29293.0000 - val_tn: 63819.0000 -
val_tp: 1227.0000 - val_binary_accuracy: 0.6839 - val_precision: 0.0402 -
val_recall: 0.6157

Epoch 14/20

221909/221909 [=====] - 1s 4us/step - loss: 3.1463e-06
- fn: 1617.0000 - fp: 75627.0000 - tn: 141477.0000 - tp: 3188.0000 -
binary_accuracy: 0.6519 - precision: 0.0404 - recall: 0.6635 - val_loss:
3.1450e-06 - val_fn: 800.0000 - val_fp: 28068.0000 - val_tn: 65044.0000 -
val_tp: 1193.0000 - val_binary_accuracy: 0.6965 - val_precision: 0.0408 -
val_recall: 0.5986

Epoch 15/20

221909/221909 [=====] - 1s 4us/step - loss: 3.1409e-06
- fn: 1642.0000 - fp: 73344.0000 - tn: 143760.0000 - tp: 3163.0000 -
binary_accuracy: 0.6621 - precision: 0.0413 - recall: 0.6583 - val_loss:
3.1374e-06 - val_fn: 826.0000 - val_fp: 27043.0000 - val_tn: 66069.0000 -
val_tp: 1167.0000 - val_binary_accuracy: 0.7070 - val_precision: 0.0414 -
val_recall: 0.5855

Epoch 16/20

221909/221909 [=====] - 1s 4us/step - loss: 3.1386e-06
- fn: 1655.0000 - fp: 72319.0000 - tn: 144785.0000 - tp: 3150.0000 -
binary_accuracy: 0.6666 - precision: 0.0417 - recall: 0.6556 - val_loss:
3.1612e-06 - val_fn: 867.0000 - val_fp: 25134.0000 - val_tn: 67978.0000 -
val_tp: 1126.0000 - val_binary_accuracy: 0.7266 - val_precision: 0.0429 -
val_recall: 0.5650

Epoch 17/20

221909/221909 [=====] - 1s 4us/step - loss: 3.1341e-06
- fn: 1625.0000 - fp: 73162.0000 - tn: 143942.0000 - tp: 3180.0000 -
binary_accuracy: 0.6630 - precision: 0.0417 - recall: 0.6618 - val_loss:
3.1424e-06 - val_fn: 728.0000 - val_fp: 30863.0000 - val_tn: 62249.0000 -
val_tp: 1265.0000 - val_binary_accuracy: 0.6678 - val_precision: 0.0394 -
val_recall: 0.6347

Epoch 18/20

221909/221909 [=====] - 1s 3us/step - loss: 3.1311e-06
- fn: 1603.0000 - fp: 74621.0000 - tn: 142483.0000 - tp: 3202.0000 -
binary_accuracy: 0.6565 - precision: 0.0411 - recall: 0.6664 - val_loss:
3.1402e-06 - val_fn: 696.0000 - val_fp: 32201.0000 - val_tn: 60911.0000 -
val_tp: 1297.0000 - val_binary_accuracy: 0.6541 - val_precision: 0.0387 -
val_recall: 0.6508

Epoch 19/20

221909/221909 [=====] - 1s 3us/step - loss: 3.1378e-06
- fn: 1638.0000 - fp: 73479.0000 - tn: 143625.0000 - tp: 3167.0000 -
binary_accuracy: 0.6615 - precision: 0.0413 - recall: 0.6591 - val_loss:
3.1380e-06 - val_fn: 649.0000 - val_fp: 34746.0000 - val_tn: 58366.0000 -
val_tp: 1344.0000 - val_binary_accuracy: 0.6278 - val_precision: 0.0372 -

```

val_recall: 0.6744
Epoch 20/20
221909/221909 [=====] - 1s 3us/step - loss: 3.1286e-06
- fn: 1592.0000 - fp: 75259.0000 - tn: 141845.0000 - tp: 3213.0000 -
binary_accuracy: 0.6537 - precision: 0.0409 - recall: 0.6687 - val_loss:
3.1560e-06 - val_fn: 804.0000 - val_fp: 27750.0000 - val_tn: 65362.0000 -
val_tp: 1189.0000 - val_binary_accuracy: 0.6998 - val_precision: 0.0411 -
val_recall: 0.5966

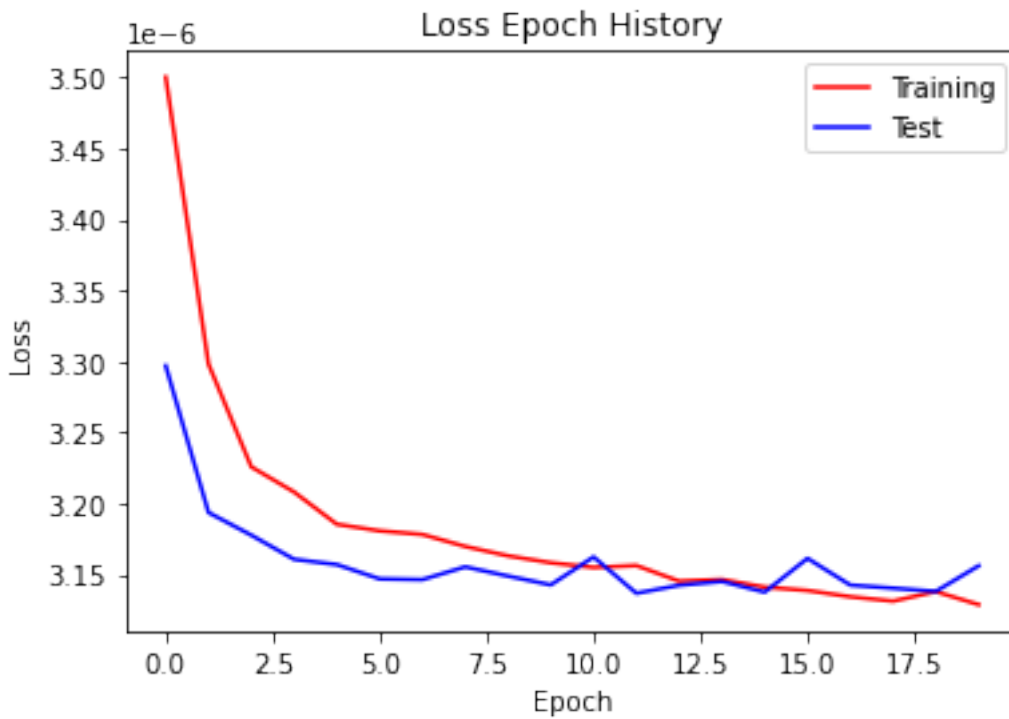
```

```
[20]: history.history.keys()
```

```
[20]: dict_keys(['val_loss', 'val_fn', 'val_fp', 'val_tn', 'val_tp',
'val_binary_accuracy', 'val_precision', 'val_recall', 'loss', 'fn', 'fp', 'tn',
'tp', 'binary_accuracy', 'precision', 'recall'])
```

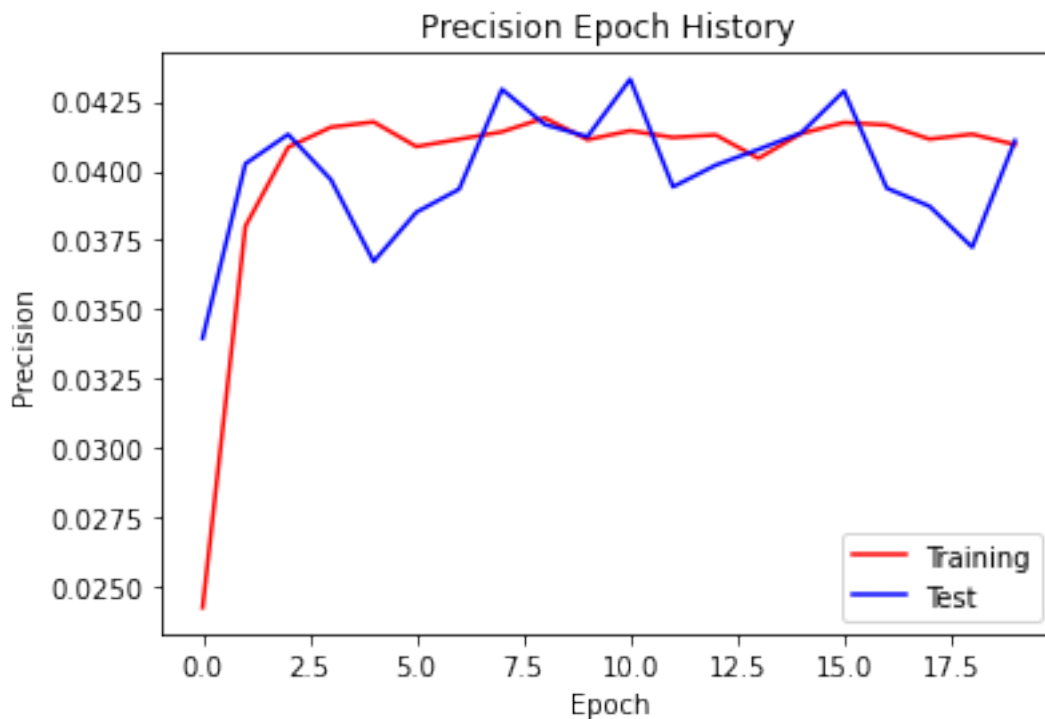
```
[21]: # Inspired by https://machinelearningmastery.com/
→,→display-deep-learning-model-training-history-in-keras/
plt.plot(history.history['loss'],'r')
plt.plot(history.history['val_loss'],'b')
plt.title('Loss Epoch History')
plt.legend(['Training','Test'])
plt.xlabel('Epoch')
plt.ylabel('Loss')
```

```
[21]: Text(0, 0.5, 'Loss')
```




```
[22]: # Inspired by https://machinelearningmastery.com/
↳,↳display-deep-learning-model-training-history-in-keras/
plt.plot(history.history['precision'],'r')
plt.plot(history.history['val_precision'],'b')
plt.legend(['Training','Test'])
plt.title('Precision Epoch History')
plt.xlabel('Epoch')
plt.ylabel('Precision')
```

```
[22]: Text(0, 0.5, 'Precision')
```



```
[23]: # Inspired by https://machinelearningmastery.com/
↳,↳display-deep-learning-model-training-history-in-keras/
plt.plot(history.history['recall'],'r')
plt.plot(history.history['val_recall'],'b')
plt.legend(['Training','Test'])
plt.title('Recall Epoch History')
plt.xlabel('Epoch')
plt.ylabel('Recall')
```

```
[23]: Text(0, 0.5, 'Recall')
```



0.6 Testing model


```
[25]: preds = model.predict(X_test)
y_pred = np.where(preds>0.5,1,0)
conf_matrix = confusion_matrix(y_test, y_pred)
conf_matrix
```

```
[25]: array([[54650, 22966],
[ 634, 1003]], dtype=int64)
```

```
[26]: conf_df = pd.DataFrame(conf_matrix, index = ['Actual Negative', 'Actual_
↪Positive'],
                        columns = ['Predict Negative', 'Predict Positive'])
sns.heatmap(conf_df, annot=True, fmt='.5g')
plt.xlabel('Predicted Label')
plt.ylabel('Actual Label')
plt.title('Confusion Matrix for Test Data')
```

UserWarning: Distributing <class 'numpy.ndarray'> object. This may take some time.

```
[26]: Text(0.5, 1.0, 'Confusion Matrix for Test Data')
```

0.7 Leveraging SHAP (SHapley Additive exPlanations) to explain test predictions

```
[145]: shap.initjs()
```

<IPython.core.display.HTML object>

```
[146]: # inspired by https://shap.readthedocs.io/en/latest/example_notebooks/
↳ tabular_examples/neural_networks/
↳ Census%20income%20classification%20with%20Keras.html
def f(X):
    return model.predict([X[:,i] for i in range(X.shape[1])]).flatten()
```

```
[147]: with open("./PROCESSED/scalers.pkl","rb") as f:
        scaler_dct = pickle.load(f)
```

```
[148]: test_df.columns
```

```
[148]: Index(['Loan Sequence Number', 'Loan Delinquency Within Year', 'Credit Score',
'MortgageInsuranceFlag', 'Units_1', 'Units_2', 'Units_3', 'Units_4',
'OccupancyStatus_I', 'OccupancyStatus_P', 'OccupancyStatus_S',
'Original Debt-to-Income (DTI) Ratio', 'Original UPB',
'Original Loan-to-Value (LTV)', 'Original Interest Rate', 'Channel_B',
```

```

'Channel_C', 'Channel_R', 'PropertyType_CO', 'PropertyType_MH',
'PropertyType_PU', 'PropertyType_SF', 'LoanPurpose_C', 'LoanPurpose_N',
'LoanPurpose_P', 'LoanTerm_360', 'LoanTerm_180', 'LoanTerm_240',
'LoanTerm_120', 'LoanTerm_300', 'OneBorrower', 'AffordableProgramFlag'],
dtype='object')

```

```
[149]: test_df[test_df['OneBorrower'] == 1].head()
```

```

[149]:  Loan Sequence Number  Loan Delinquency Within Year  Credit Score  \
1          F115Q1238542                                0.0      -0.415634
2          F119Q2306763                                0.0      -0.231641
4          F116Q1069514                                0.0      -0.605955
5          F111Q3042379                                0.0       1.858747
8          F117Q4058814                                0.0       0.899534

      MortgageInsuranceFlag  Units_1  Units_2  Units_3  Units_4  \
1                          1         1         0         0         0
2                          0         1         0         0         0
4                          0         1         0         0         0
5                          0         1         0         0         0
8                          0         1         0         0         0

      OccupancyStatus_I  OccupancyStatus_P  ...  LoanPurpose_C  LoanPurpose_N  \
1                      0                  1  ...              0              0
2                      0                  1  ...              1              0
4                      0                  1  ...              0              0
5                      0                  1  ...              0              0
8                      1                  0  ...              1              0

      LoanPurpose_P  LoanTerm_360  LoanTerm_180  LoanTerm_240  LoanTerm_120  \
1                  1             1             0             0             0
2                  0             1             0             0             0
4                  1             1             0             0             0
5                  1             1             0             0             0
8                  0             1             0             0             0

      LoanTerm_300  OneBorrower  AffordableProgramFlag
1                 0            1                      0
2                 0            1                      0
4                 0            1                      0
5                 0            1                      0
8                 0            1                      0

[5 rows x 32 columns]

```

```
[150]: test_df[test_df['Credit Score'] < -5].head()
```

```
[150]:      Loan Sequence Number  Loan Delinquency Within Year  Credit Score  \
16569      F110Q1242556                                0.0      -5.199338

      MortgageInsuranceFlag  Units_1  Units_2  Units_3  Units_4  \
16569                      0         1         0         0         0

      OccupancyStatus_I  OccupancyStatus_P  ...  LoanPurpose_C  \
16569                  0                  1  ...              1

      LoanPurpose_N  LoanPurpose_P  LoanTerm_360  LoanTerm_180  LoanTerm_240  \
16569              0              0           0           1           0

      LoanTerm_120  LoanTerm_300  OneBorrower  AffordableProgramFlag
16569             0             0           0                      0

[1 rows x 32 columns]
```

```
[151]: # Converting features back to original scale for display purposes
def get_original(col,val):
    if col in scaler_dct:
        scaler = scaler_dct[col]
        nval = scaler.inverse_transform(np.array([[val]]))[0][0]
    else:
        nval = val
    return nval
```

```
[152]: rec = X_test[4]
shap_values = explainer.shap_values(rec, nsamples=2500)
feature_vals = [get_original(col_name,value) for col_name,value in
↳zip(keep_cols[2:],rec)]
feature_vals = np.array(feature_vals)
shap.force_plot(explainer.expected_value, shap_values[0],
                features=feature_vals, feature_names=keep_cols[2:])
```

```
[152]: <shap.plots._force.AdditiveForceVisualizer at 0x1688249cf98>
```

```
[153]: shap_values100 = explainer.shap_values(X_test[100:110], nsamples=500)
```

```
0%|          | 0/10 [00:00<?, ?it/s]
```

```
[154]: shap.force_plot(explainer.expected_value, shap_values100[0], X_test[100:110],
↳feature_names=keep_cols[2:])
```

```
[154]: <shap.plots._force.AdditiveForceArrayVisualizer at 0x16892859320>
```

```
[ ]:
```